

Presentation

Gilbert: A sparse linear algebra environment

Till Rohrmann

`till.rohrmann@campus.tu-berlin.de`

Technische Universität Berlin

April 2, 2014

Table of Contents

- 1 Motivation
- 2 Gilbert
- 3 Approach
- 4 Current state
- 5 Outlook
- 6 Related Work


Big data analytics

- Gathered data grows exponentially
- Obtain new insights
- Analytic methods have to scale up ⇒ Parallelization
- Methods developed within linear algebra systems
- Explicit parallelization tedious and error-prone



Distributed computing and data analytics

- Experts familiar with both domains countable
- Laborious to become acquainted with new domain
- Huge existing code base
- Can't we bring both worlds together?
- Solution: Gilbert



Gilbert

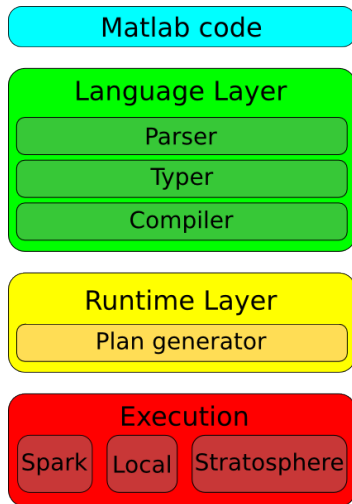
- Distributed sparse linear algebra system
- Matlab frontend for distributed computing frameworks
- Allows to almost seamlessly move from local execution to distributed execution in a heterogenous environment
- Enable machine learning and data analytics algorithms to be run distributedly

Approach

- 1 Compiling Matlab code into intermediate representation
- 2 Apply optimizations indenpendently of runtime specific system
- 3 Compiling intermediate representation into runtime specific format



System architecture



Language

- Matlab-like language
- Support of basic linear algebra operations
- Some built-in functions, repmat, linspace, pdist2
- Loop support with static and dynamic termination criterion
- Language expressive enough to support variety of algorithms, Pagerank, K-means, NNMF

```
1  A'*B
```

```
3  f = @(x) x.^2.0
```

```
5  eps = 0.1
```

```
7  c = @(p,c) norm(p-c,2) < eps
```

```
9  fixpoint(1/2, f, 10, c)
```


Compilation: Intermediate format

- Scala's combinator parsing tool box
- Powerful enough for our language
- Matlab is dynamically typed
- Execution on Stratosphere requires type knowledge at compilation time
- Hindley-Milner type inference algorithm to infer types and dimensions

Example: Parsing and typing

Input

```
A = ones(2,10);  
B = eye(10,3);  
A*B
```

Parsed and typed

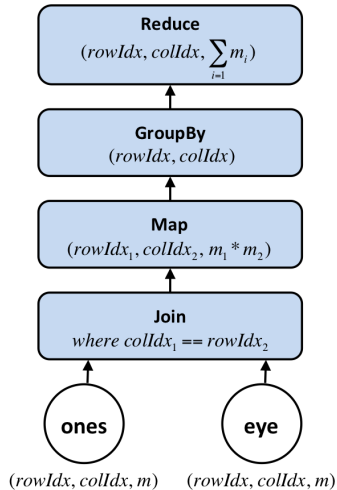
```
A = ones(2,10): MatrixType[ Double, 2, 10];  
B = eye(10,3): MatrixType[ Double, 10, 3];  
A*B: MatrixType[ Double, 2, 3]
```

Example: Intermediate code

Compiled

```
MatrixMult(  
  ones(  
    IntValue(2),  
    IntValue(10)  
  ),  
  eye(  
    IntValue(10),  
    IntValue(3)  
  )  
)
```

Example: Stratosphere execution plan



Current state

- Language layer implemented and working
- Runtime layer: Execution on Stratosphere with iteration and convergence support
- Implemented algorithms: PageRank, NNMF and K-means



Live demonstration



Outlook

- Typing system with constraints, similar to Haskell typing system with type class support
- System evaluation: Runtime, scalability
- Comparison to specialized algorithms
- Optimizations for the intermediate representation



Related Work

- SystemML [2]
 - Higher-level language for linear algebra
 - Compiled to MapReduce jobs
- Apache Mahout [1]
 - Specialized implementations of algorithms of various kind
 - No native linear algebra support
- Pegasus [3]
 - Generalized iterative matrix vector multiplication
- Spark [4]
 - MapReduce extension



Bibliography I



Apache. *Apache Mahout*. Cited January 9th 2014. 2011.



A. Ghoting et al. “SystemML: Declarative machine learning on MapReduce”. In: *Data Engineering (ICDE), 2011 IEEE 27th International Conference on*. IEEE. 2011, pp. 231–242.



U Kang, C. E. Tsourakakis, and C. Faloutsos. “Pegasus: A peta-scale graph mining system implementation and observations”. In: *Data Mining, 2009. ICDM'09. Ninth IEEE International Conference on*. IEEE. 2009, pp. 229–238.



M. Zaharia et al. “Spark: cluster computing with working sets”. In: *Proceedings of the 2nd USENIX conference on Hot topics in cloud computing*. 2010, pp. 10–10.

Replication based matrix multiplication

$$\begin{bmatrix} A_{0,0} & A_{0,1} \\ A_{1,0} & A_{1,1} \end{bmatrix} * \begin{bmatrix} B_0 \\ B_1 \end{bmatrix} = \begin{bmatrix} C_0 \\ C_1 \end{bmatrix}$$


Node 1

$$\begin{bmatrix} A_{0,0} & A_{0,1} \end{bmatrix} * \begin{bmatrix} B_0 \\ B_1 \end{bmatrix} = \begin{bmatrix} C_0 \end{bmatrix}$$

Node 2

$$\begin{bmatrix} A_{1,0} & A_{1,1} \end{bmatrix} * \begin{bmatrix} B_0 \\ B_1 \end{bmatrix} = \begin{bmatrix} C_1 \end{bmatrix}$$

Cross product based matrix multiplication



Communication costs

- Matrix A blocked with $M \times K$ blocks and B blocked with $K \times N$ blocks
- Network and IO costs are dominant
- $costs_{RMM} \in O(network(|A| \cdot N + |B| \cdot M) + io(|A| + |B| + |C|))$
- $costs_{CPMM} \in O(network(|A| + |B| + r \cdot |C|) + io(|A| + |B| + |C|))$
with r being the number of reducer

Stratosphere execution plan

- Assuming row-wise partitioning of matrix A
- Execution plans for RMM and CPMM identical
- Differ only in chosen strategy for join operator
- RMM: Hybrid-hash join
- CPMM: Sort-merge join
- Stratosphere's optimizer chooses right plan

